

HPC

“Evaluate performance enhancement of parallel Quicksort Algorithm using MPI”

✓ 1. What is the project about?

This project focuses on improving the performance of the Quicksort algorithm by making it parallel using MPI (Message Passing Interface).

Normally, Quicksort runs on a single core and sorts data sequentially, but as data sizes grow and multi-core systems become common, this approach is slow. So, we parallelize Quicksort across multiple processors to speed up the sorting process, especially for large datasets.

We tested this parallel implementation and compared it with the traditional (sequential) version to evaluate improvements in:

Execution time

Speedup (how much faster)

Efficiency (how well processors are utilized)

✓ 2. Why we chose this project (Relevance + Motivation)

Sorting is a fundamental operation in computing, used everywhere from databases to simulations.

Quicksort is a popular sorting algorithm because it's fast on average.

But on large datasets, sequential Quicksort wastes modern hardware resources like multi-core CPUs and distributed clusters.

MPI allows programs to run across multiple processors, even across machines.

By using MPI, we can divide the data, sort in parallel, and merge efficiently, reducing total time.

This is highly useful in:

High-Performance Computing (HPC)

Big data analytics

Scientific simulations

✓ 3. Problem statement

The problem with traditional Quicksort:

Runs on one processor → slow for huge datasets.

Cannot take advantage of multi-core or distributed systems.

Sequential approach becomes a bottleneck.

Challenges in parallelizing Quicksort:

Splitting data efficiently across processes.

Managing inter-process communication.

Ensuring load balancing (each processor gets roughly equal work).

This project solves these challenges using MPI.

✓ 4. Scope of the project

Implement parallel Quicksort using MPI.

Run experiments to compare parallel vs. sequential performance.

Measure:

Execution time

Speedup

Efficiency across datasets and processor counts

Study trade-offs:

Communication overhead

Load balancing

✓ 5. How does the parallel Quicksort work?

Here's an easy-to-understand explanation:

Initialization

Start MPI using `MPI_Init()`.

Find number of processes (`MPI_Comm_size`) and process rank (`MPI_Comm_rank`).

Data distribution

Root process (rank 0) generates a random dataset.

Split the dataset equally among all processes using `MPI_Scatter`.

Local sorting

Each process runs Quicksort on its local data chunk.

This happens in parallel, cutting down the overall sorting time.

Data merging

After local sorts, processes combine their sorted chunks.

Pairwise merge (e.g., rank 0 ↔ rank 1, rank 2 ↔ rank 3, etc.).

Use a merge function and keep appropriate halves.

Final gathering

Final sorted data is gathered at root using MPI_Gather or MPI_Bcast.

Termination

Close MPI with MPI_Finalize().

✓ 6. Tools and software

Hardware: Multi-core PC or cluster

Software:

Python

mpi4py library

OpenMPI

Jupyter Notebook or Google Colab

✓ 7. Code highlights (how to explain code if examiner asks)

mpi4py → Provides Python bindings for MPI.

quicksort() function → Standard recursive Quicksort, uses a random pivot.

Scatter → Distributes data chunks.

Each process → Sorts its local chunk.

Gather → Collect sorted lists.

Final sort → Final merge and Quicksort at root.

✓ 8. Results and discussion

Large datasets (>100,000 elements):

Significant speedup with MPI parallel Quicksort.

Small datasets:

Overhead of communication reduces gains.

More processors:

Speedup improves initially, but after a point, overhead increases → diminishing returns.

Key takeaway:

Optimal performance depends on dataset size + number of processors.

✓ 9. Conclusion

Parallel Quicksort using MPI improves performance over sequential Quicksort.

Works well on large datasets.

Best performance when processor count and data size are balanced.

Future improvements could include:

Optimized merging

Dynamic load balancing

Use on distributed clusters

✓ 10. Real-world applications

Sorting in distributed databases

Scientific simulations

Big data processing

Sorting logs or sensor data on HPC clusters

🌟 1-minute summary you can memorize

“Our project implemented a parallel version of the Quicksort algorithm using MPI to improve sorting speed on large datasets. We divided data among multiple processors, sorted each part locally, and merged results. We measured execution time, speedup, and efficiency, finding that parallel Quicksort outperforms the sequential version, especially on large datasets. The project highlights the benefits and trade-offs of parallel computing, with real-world uses in HPC, data analytics, and scientific applications.”

🔥 Likely examiner viva questions + answers

Question Answer

Why Quicksort? It's efficient on average and widely used, but its recursive design makes it a good candidate for parallelization.

Why MPI? MPI allows communication between multiple processes, even on different machines, enabling distributed parallelism.

Challenges in parallelizing Quicksort? Data division, communication cost, and load balancing.

What metrics did you use? Execution time, speedup, efficiency.

What are the limitations? Communication overhead, diminishing returns when too many processors are used.